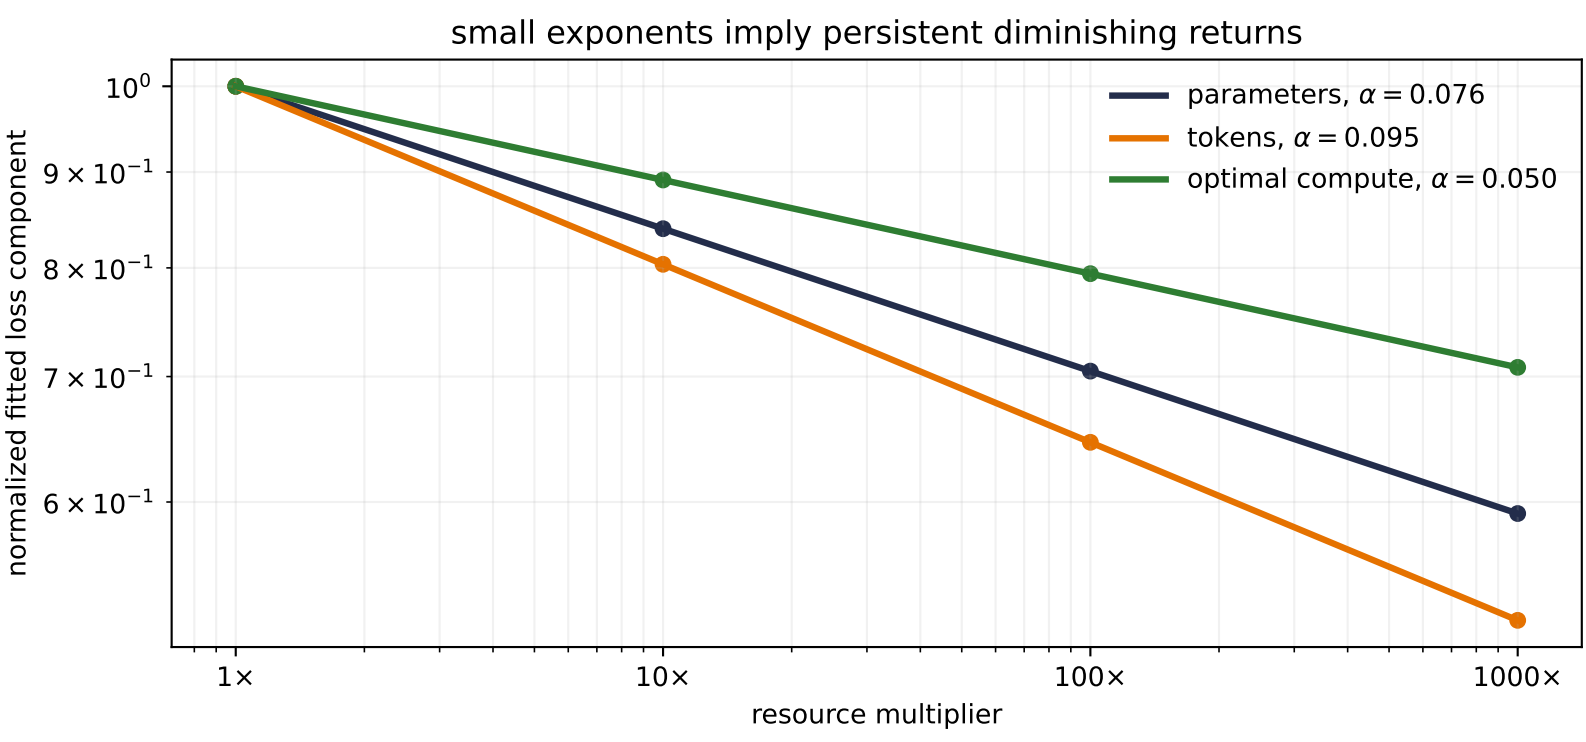
small exponents imply persistent diminishing returns
normalized fitted loss component
10^0
9×10^{-1}
8×10^{-1}
7×10^{-1}
6×10^{-1}
parameters, $\alpha = 0.076$
tokens, $\alpha = 0.095$
optimal compute, $\alpha = 0.050$
1×
10×
100×
1000×
resource multiplier