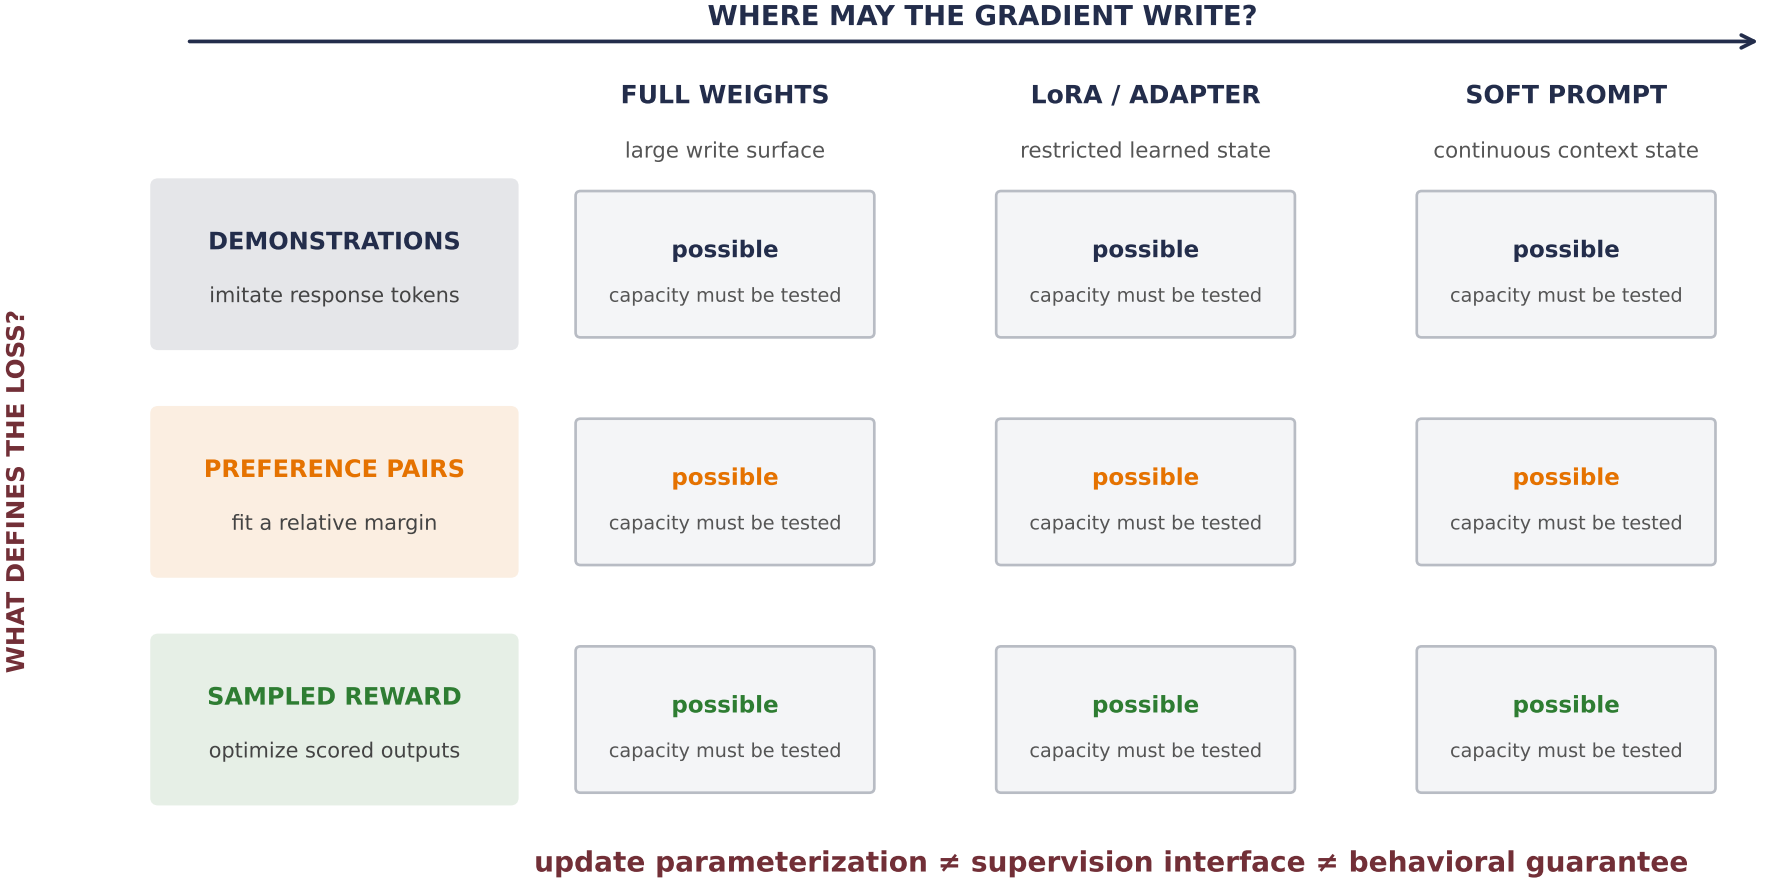
WHERE MAY THE GRADIENT WRITE?
WHAT DEFINES THE LOSS?
FULL WEIGHTS
large write surface
LoRA / ADAPTER
restricted learned state
SOFT PROMPT
continuous context state
DEMONSTRATIONS
imitate response tokens
PREFERENCE PAIRS
fit a relative margin
SAMPLED REWARD
optimize scored outputs
possible
capacity must be tested
update parameterization ≠ supervision interface ≠ behavioral guarantee