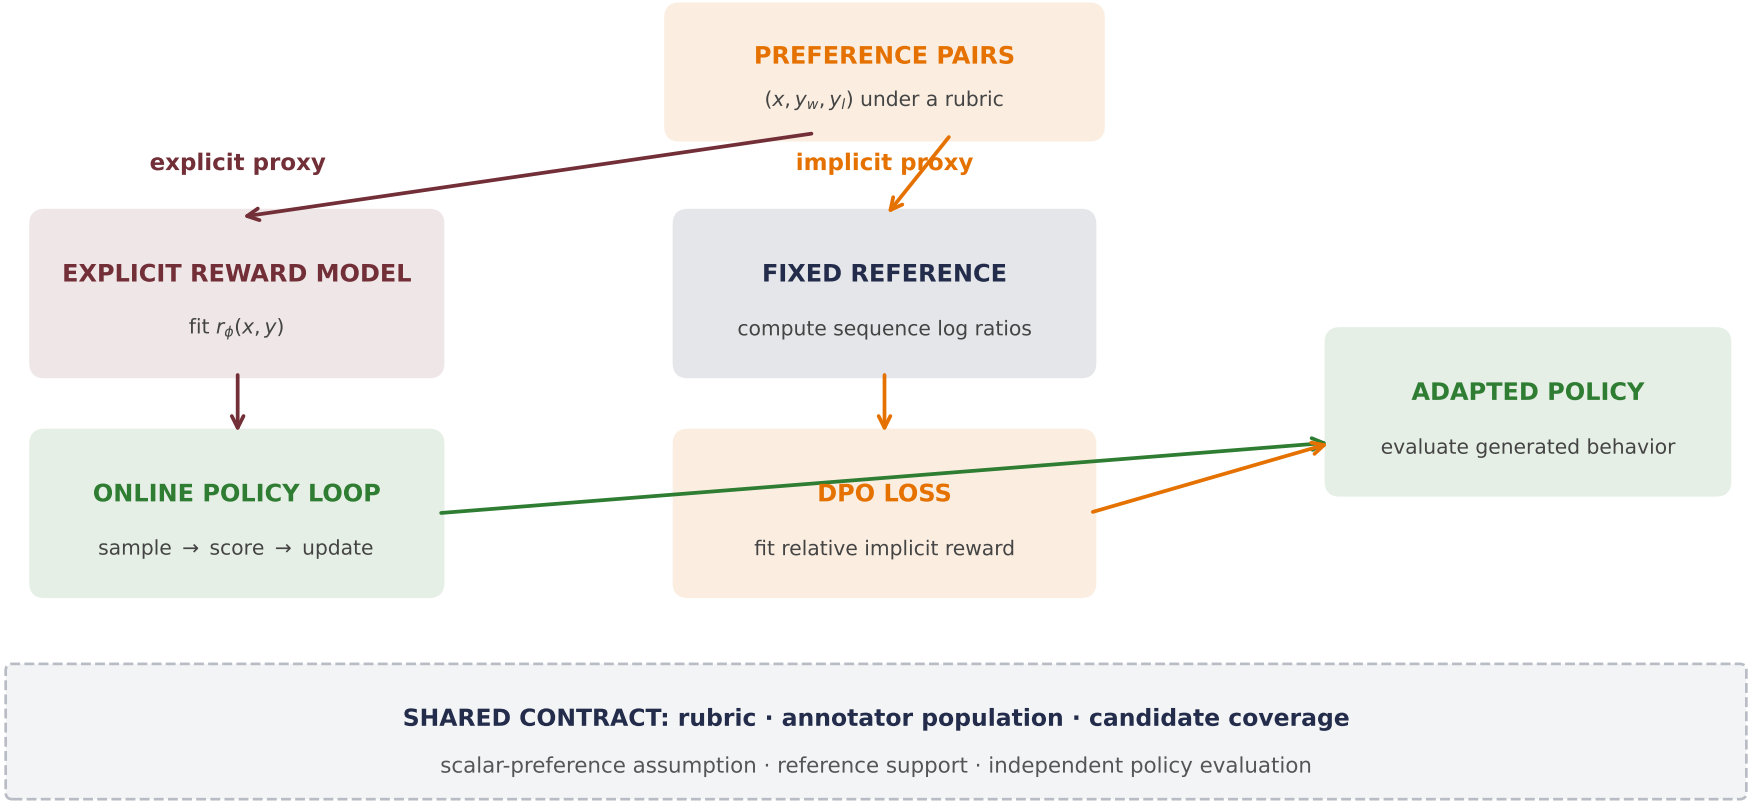
PREFERENCE PAIRS
(x, y_w, y_l) under a rubric
explicit proxy
implicit proxy
EXPLICIT REWARD MODEL
fit $r_\phi(x, y)$
FIXED REFERENCE
compute sequence log ratios
ADAPTED POLICY
evaluate generated behavior
ONLINE POLICY LOOP
sample → score → update
DPO LOSS
fit relative implicit reward
SHARED CONTRACT: rubric · annotator population · candidate coverage
scalar-preference assumption · reference support · independent policy evaluation