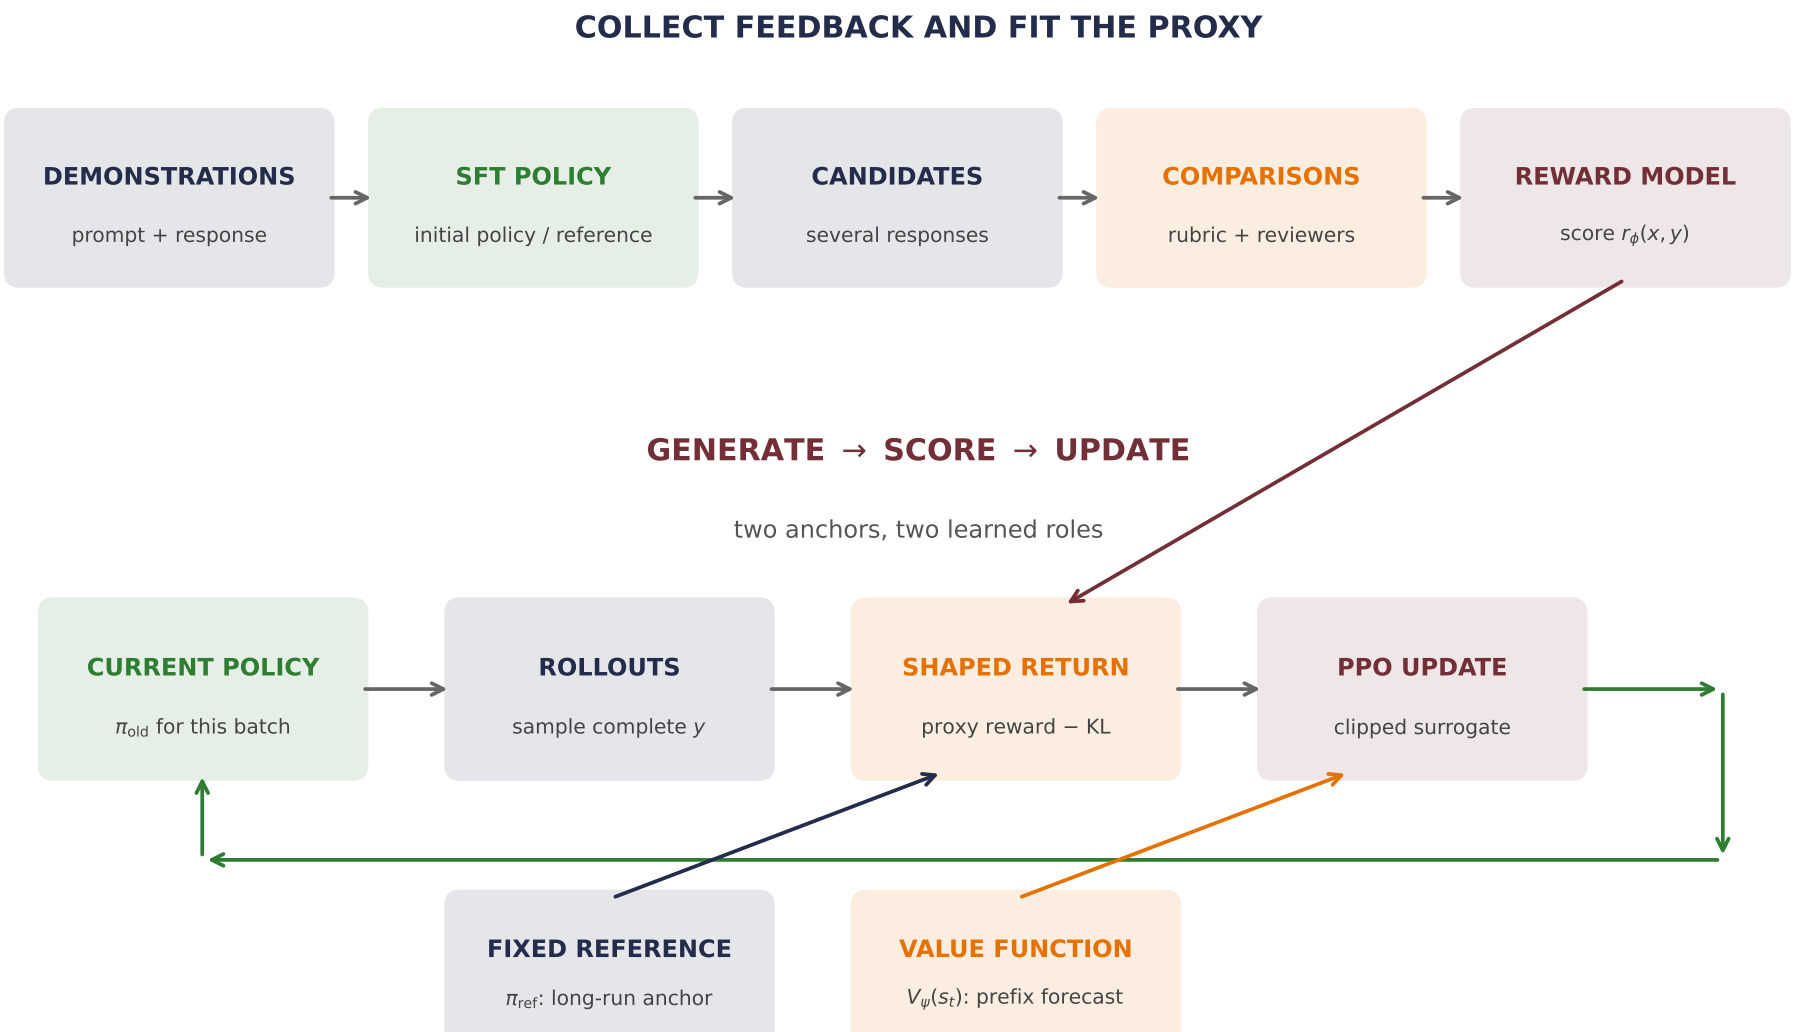
COLLECT FEEDBACK AND FIT THE PROXY
DEMONSTRATIONS
prompt + response
SFT POLICY
initial policy / reference
CANDIDATES
several responses
COMPARISONS
rubric + reviewers
REWARD MODEL
score $r_\phi(x, y)$
GENERATE → SCORE → UPDATE
two anchors, two learned roles
CURRENT POLICY
π_{old} for this batch
ROLLOUTS
sample complete y
SHAPED RETURN
proxy reward − KL
PPO UPDATE
clipped surrogate
FIXED REFERENCE
π_{ref}: long-run anchor
VALUE FUNCTION
$V_\psi(s_t)$: prefix forecast