

the approximate posterior can miss evidence = ELBO + posterior KL

generation and inference

PRIOR

$$z \sim p(z)$$

DECODER

$$p_{\theta}(x | z)$$

sample path

INPUT

$$x$$

ENCODER

$$q_{\phi}(z | x)$$

training-time approximation

